



Introduction to Modern System Design

Get an overview of the topics we'll cover in this course.

We'll cover the following

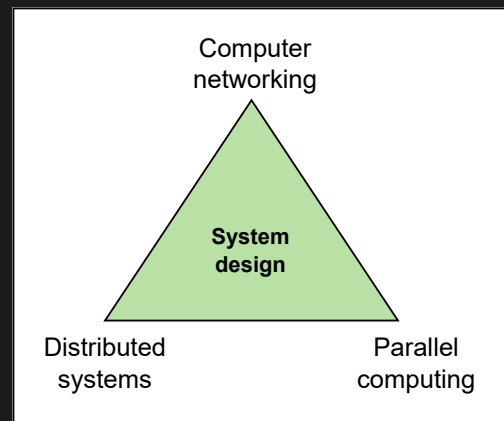


- What is system design?
 - Modern system design using building blocks
 - About this course
 - Who should take this course?
 - Prerequisites for this course

What is system design?

System design is the process of defining components and their integration, APIs, and data models to build large-scale systems that meet a specified set of functional and non-functional requirements.

System design uses the concepts of computer networking, parallel computing, and distributed systems to craft systems that scale well and are performant. Distributed systems scale well by nature. However, distributed systems are inherently complex. The discipline of system design helps us tame this complexity and get the work done.



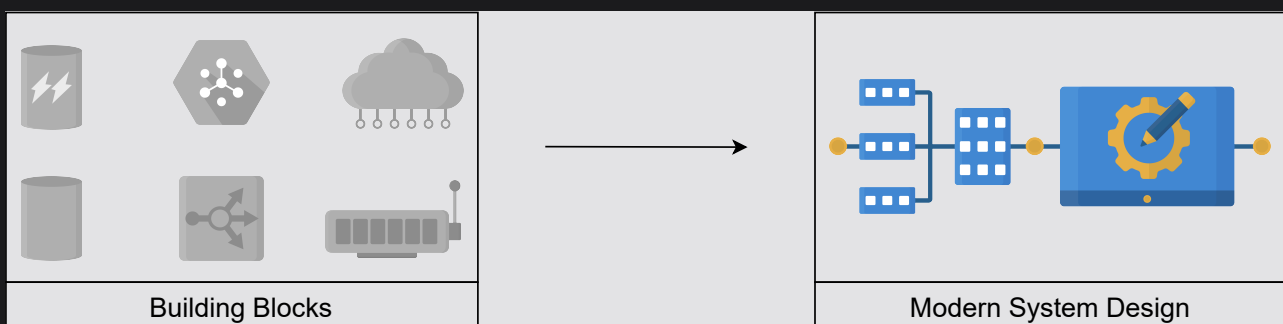
System design aims to build systems that are reliable, effective, and maintainable, among other characteristics.

- **Reliable systems** handle faults, failures, and errors.
- **Effective systems** meet all user needs and business requirements.
- **Maintainable systems** are flexible and easy to scale up or down. The ability to add new features also comes under the umbrella of maintainability.

Modern system design using building blocks

We have separated out commonly-used design elements, such as load balancers, as the basic building blocks for high-level system design. This serves two purposes. First, it allows us to discuss all the building blocks in detail and discuss their interesting mini-design problems. Second, when we tackle a design problem, we can concentrate on problem-specific aspects, mention the building block we'll use, and how we'll use it. This helps us remove duplicate discussions of commonly-occurring design elements.

We have identified sixteen building blocks that are crucial in designing modern systems.



Modern systems are designed using fundamental building blocks

About this course

This course is about designing systems that scale with increasing users, remain available even under different faults, and meet functional goals with good

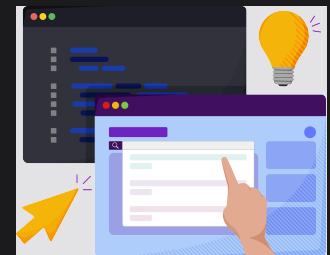


performance. Real-world system building is an iterative process where we start with a reasonably good design, measure how it performs, and improve the design in the next iteration.



The focus of this course is to immerse ourselves into carefully-selected system design endeavors to enable ourselves to tackle any novel design problem, be it in a systems design interview or a task at the office. This course aims to teach concepts instead of giving out boilerplate designs. Some gaps that this course aims to fill are listed below.

A fresh look at system design: Many system design courses provide a formula to attack a specific problem. This might seem attractive in a high-stress situation like an interview, but it might encourage memorizing a design solution instead of actually understanding the problem and devising an appropriate solution. If system design were that formulaic, then we probably wouldn't need people for system designing. System design is as much an art as it is a science, and attacking a design problem from the first principles gives a fresh feel to it.



Going deep and broad: We tackle some traditional problems, but with added in-depth discussions on them. We give proper rationale for why we use some components despite their tradeoffs. For example, we explain why we use a particular database, a caching system, or a load balancing technique in a design.



We address some new design problems as well that touch upon not only scalability but also availability, maintainability, consistency, and fault-tolerance. Collectively, traditional and new problems cover all aspects of modern system design activity. Our hope is that



this course prepares learners to effectively tackle any new design problem they encounter.



Real systems are complex and, often, we might need to make appropriate assumptions to properly scope a problem. We cover problems in more detail to properly grasp the real-world systems.

Iterative process: Systems, in reality, improve over iterations. We often start with something simple, but when



educative

problems, we make one design, identify bottlenecks, and improve on it. Working under time constraints might not permit iterations on the design. However, we still recommend two iterations—first, where we do our best to come up with a design (that takes about 80 percent of our time), and a second iteration for improvements. Another choice is to change things as we figure out new insights. Inevitably, we discover new details as we spend more time working with a problem.

Interactive learning: We provide ample opportunities to get experience with system design. Some design problems guide learners through many steps to design a system. We also have a few examples where the learner designs the full system end-to-end without any guided steps. We reinforce the important concepts by testing learners with questions and quizzes.



Who should take this course?

System design is for any software engineer who wants to advance in their career.





- **Interview preparation:** Lately, system design is becoming an important part of software development interviews. This course helps software engineers prepare for interviews. We have an elaborate guide on preparation for a system design interview in the second chapter of this course for learners who are interested.
- **Software developers:** System design is primarily for back-end developers who aim to become principal engineers or solution architects. It's because these engineers handle actual user data. Once the data is submitted to back-end systems by the frontend, effective handling of data makes the overall application successful. Having said that, full stack or front-end developers may also want to learn system design to improve their work. At the same time, **support engineers** (also called **SREs**) who work on-call in the production environment have to deal with all sorts of problems daily. Therefore, system design concepts enable SREs to efficiently find the root causes of complex problems.
- **Product/project managers:** A big challenge in product/project management is to build systems that scale well and perform effectively over time. Managers that are aware of system design can steer the design of large-scale performant systems. Therefore, it is imperative for product/project managers to understand system design concepts to lead the design and development of successful applications.
- **System design learners:** System design is an interesting subject, and people in tech domains can greatly benefit from learning system design. This course helps a learner understand how giant tech companies design and build successful applications from scratch and improve them over time. Other learners may want to develop their idea into a large application by learning from this course.

Prerequisites for this course

We assume that you know the fundamental concepts of a distributed system.



System design borrows many great concepts from distributed systems. We have an [excellent course](#) for learners who are interested in revisiting and reviewing their concepts on distributed systems. Apart from distributed systems, some basic concepts on [computer networking](#) and [operating systems](#) are also helpful before taking this course.



Completed

Next →

Course Structure for Modern System ...

